

CSC 480/580 Principles of Machine Learning

## 02 Limits of Learning

**Chicheng Zhang**

**Department of Computer Science**




# Motivation

- Supervised learning is a general & useful framework
- Understand when supervised learning will and will not work

# Optimal classification with known $D$

- Suppose:  $I(A) = 1$  if  $A$  happens, and  $= 0$  otherwise
  - Binary classification, 0-1 loss  $\ell(y, \hat{y}) = I(y \neq \hat{y})$
  - $D$  is *known*: for every  $(x, y)$ ,  $P_D(x, y)$  is known to us
- What is the  $f$  that has the smallest *generalization error*  
 $L_D(f) = \mathbb{E}_{(x,y) \sim D} I(y \neq f(x))$ ?
- Note (alternative expression) :  $L_D(f) = P_{(x,y) \sim D} (y \neq f(x))$



Generalization error:  $L_D(f) = \mathbb{E}_{(x,y) \sim D} I(y \neq f(x))$

# Simple case: discrete domain $\mathcal{X}$

$P_D(x, y)$	$x = 1$	$x = 2$	$x = 3$
$y = -1$	0.2	0.2	0.15
$y = +1$	0.1	0.3	0.05

Which classifier is better?

- $f_1(1) = -1, f_1(2) = -1, f_1(3) = -1 \Rightarrow L_D(f_1) = 0.1 + 0.3 + 0.05$
- $f_2(1) = -1, f_2(2) = +1, f_2(3) = -1 \Rightarrow L_D(f_2) = 0.1 + 0.2 + 0.05$

Is this the best classifier? Why?

- For any  $x$ , should choose  $y$  that has higher value of  $P_D(x, y)$
- $f^*(1) = -1, f^*(2) = +1, f^*(3) = -1$

# Bayes optimal classifier

**Theorem**  $f_{BO}$  achieves the smallest generalization error among all classifiers.

$$f_{BO}(x) = \arg \max_{y \in \mathcal{Y}} P_D(X = x, Y = y) = \arg \max_{y \in \mathcal{Y}} P_D(Y = y | X = x), \forall x \in \mathcal{X}$$

**Example** Iris dataset classification:



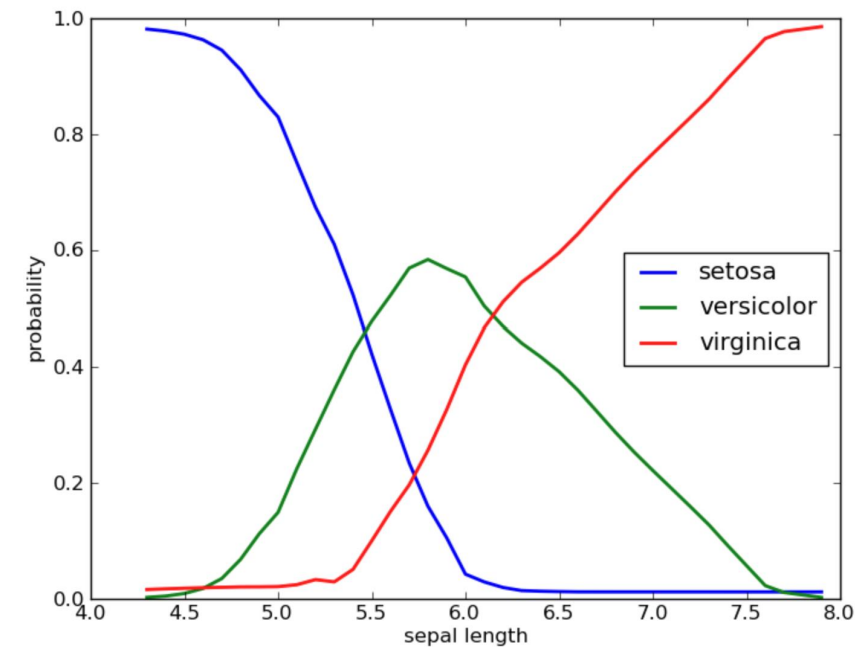
Iris Setosa



Iris Versicolor



Iris Virginica



# Proof of theorem

**Step 1** consider accuracy,

- $A_D(f) = 1 - L_D(f) = P_D(Y = f(X)) = \sum_x P_D(X = x, Y = f(x))$
- Suffices to show  $f_{BO}$  has the highest accuracy

**Step 2** comparison,

$$A_D(f_{BO}) - A_D(f) = \sum_x P_D(X = x, Y = f_{BO}(x)) - P_D(X = x, Y = f(x)) \geq 0$$

$$f_{BO}(x) = \arg \max_{y \in \mathcal{Y}} P_D(X = x, Y = y)$$

## Remarks



- Similar reasoning can be used to prove the theorem with continuous domain  $\mathcal{X}$  (sum  $\rightarrow$  integral)
- This just shows deterministic classifier, can be extended to show BO is 0-1 optimal for all classifiers

# Bayes error rate: alternative form

$$\begin{aligned} L_D(f_{BD}) &= P_D(Y \neq f_{BD}(X)) \\ &= \sum_x P_D(Y \neq f_{BD}(x) \mid X = x) P_D(X = x) \\ &= \sum_x (1 - P_D(Y = f_{BD}(x) \mid X = x)) P_D(X = x) \\ &= \sum_x \left( 1 - \max_y P_D(Y = y \mid X = x) \right) P_D(X = x) \\ &= E \left[ 1 - \max_y P_D(Y = y \mid X) \right] \end{aligned}$$

- Special case: binary classification




- $L_D(f_{BD}) = \sum_x P_D(Y \neq f_{BD}(x), X = x)$   
 $= \sum_x \min( P_D(Y = +1, X = x), P_D(Y = -1, X = x) )$



# When is the Bayes error rate nonzero?


$$L_D(f_{BO}) = \sum_x \min( P_D(Y = +1, X = x), P_D(Y = -1, X = x) )$$

- Limited feature representation
- Noise in the training data
  - Feature noise
  - Label noise – e.g. typo transcribing reviews
  - Sensor failure
  - Typo in reviews for sentiment classification
- May not have a single “correct” label






# Overfitting vs Underfitting




# Overfitting vs Underfitting

- Q: should I train a shallow or deep decision tree?

- Shallow tree:




- Deep tree:



- Underfitting: have the opportunity to learn something but didn't
- Overfitting: pay too much attention to idiosyncrasies to training data, and do not generalize well
- A model that neither overfits nor underfits is expected to do best

# Unbiased model evaluation using test data

- Your boss says: I will allow your recommendation system to run on our website only if the error is  $\leq 10\%$ !
- How to prove it?
- Idea: reserve some data as test data for evaluating predictors




- $L_{\text{test}}(\hat{f}) = \frac{1}{|S_{\text{test}}|} \sum_{(x,y) \in S_{\text{test}}} I(y \neq \hat{f}(x))$
- Law of large numbers  $\Rightarrow L_{\text{test}}(\hat{f}) \rightarrow L_D(\hat{f})$

# Law of large numbers (LLN)


- Suppose  $v_1, \dots, v_n$  are IID (independent & identically distributed) random variables, the sample average  $\bar{v} = \frac{1}{n} \sum_{i=1}^n v_i$  converges to  $E[v_1]$  as  $n \rightarrow \infty$

- Useful in e.g. election poll
- Foundation of statistics

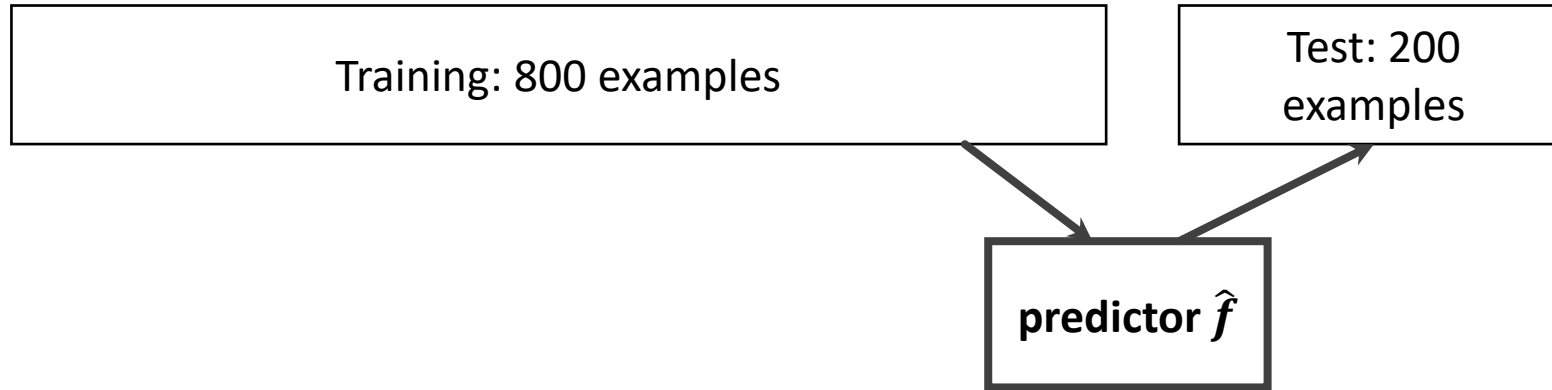


- Can we apply LLN to conclude that

- $$L_{\text{train}}(\hat{f}) = \frac{1}{|S_{\text{train}}|} \sum_{(x,y) \in S_{\text{train}}} I(y \neq \hat{f}(x)) \rightarrow L_D(\hat{f}) \text{ as } |S_{\text{train}}| \rightarrow \infty?$$



# Never touch your test data!



- More precisely: test data should only be used once, only for final evaluation
- If  $\hat{f}$  depends on test examples,  $L_{\text{test}}(\hat{f})$  may no longer estimate  $L_D(\hat{f})$  accurately
- Be mindful about indirect dependence as well:
  - adaptive data analysis – choose a new learning algorithm, after seeing that the previous algorithm produces a high-test-error model

# Case Study: MNIST Dataset

All publications use standard train/test split

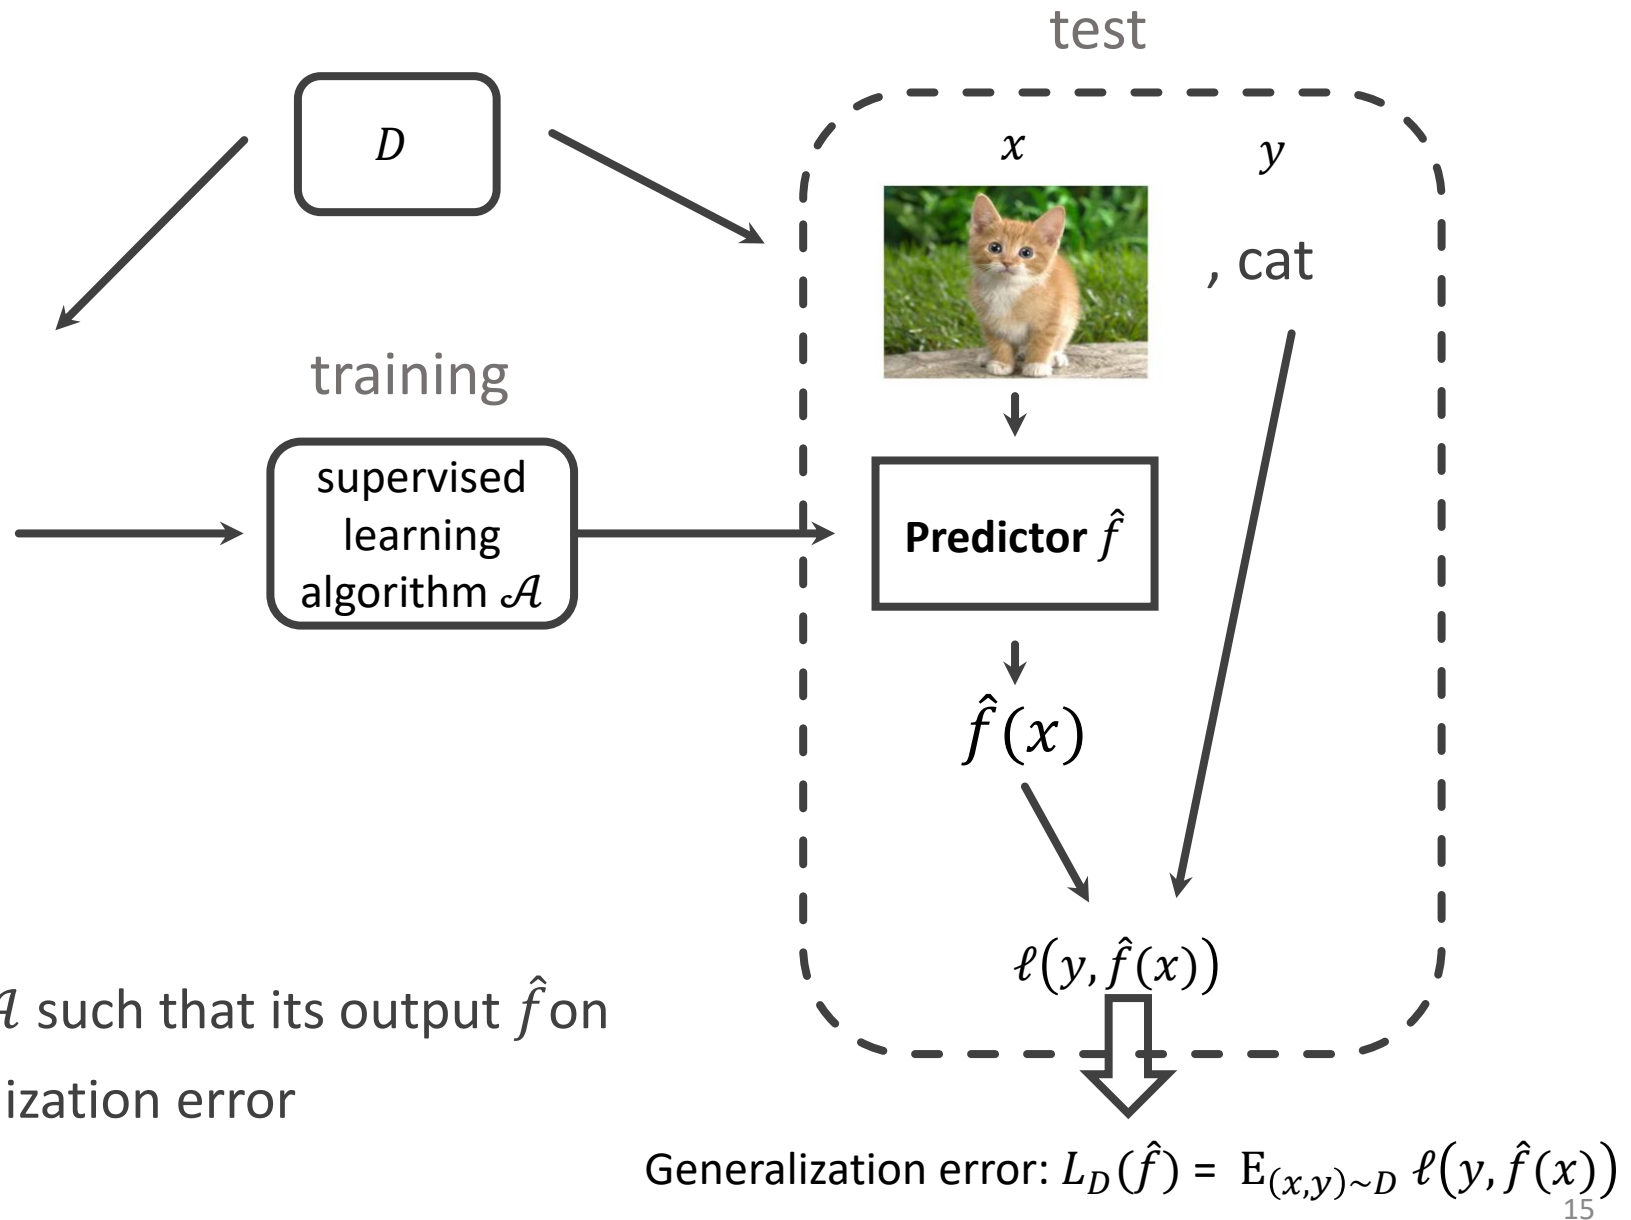
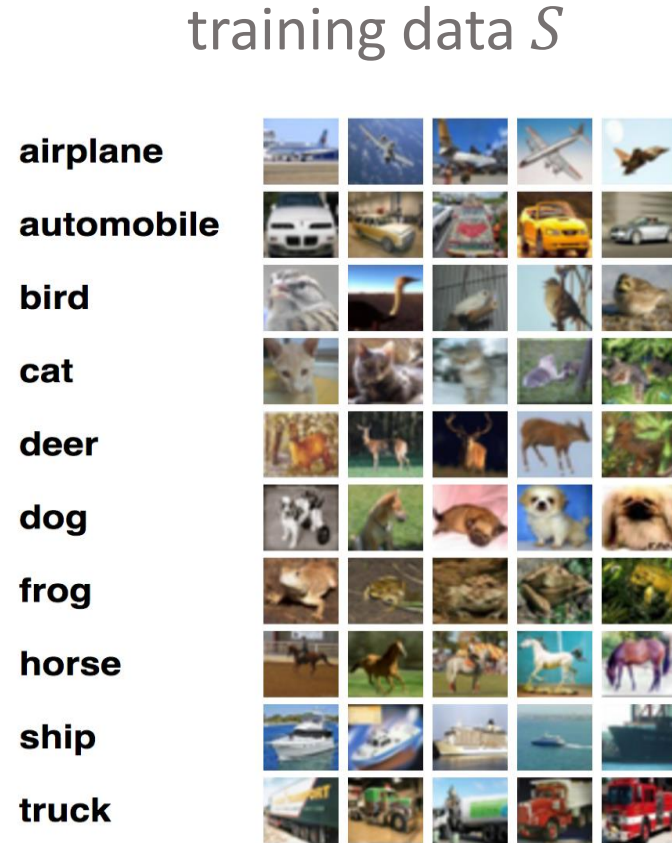
Hundreds of publications compare to each other



Type	Classifier	Distortion	Preprocessing	Error rate (%)
Linear classifier	Pairwise linear classifier	None	Deskewing	7.6 <sup>[10]</sup>
Decision stream with Extremely randomized trees	Single model (depth > 400 levels)	None	None	2.7 <sup>[28]</sup>
K-Nearest Neighbors	K-NN with rigid transformations	None	None	0.96 <sup>[29]</sup>
K-Nearest Neighbors	K-NN with non-linear deformation (P2DHMDM)	None	Shiftable edges	0.52 <sup>[30]</sup>
Boosted Stumps	Product of stumps on Haar features	None	Haar features	0.87 <sup>[31]</sup>
Non-linear classifier	40 PCA + quadratic classifier	None	None	3.3 <sup>[10]</sup>
Random Forest	Fast Unified Random Forests for Survival, Regression, and Classification (RF-SRC) <sup>[32]</sup>	None	Simple statistical pixel importance	2.8 <sup>[33]</sup>
Support-vector machine (SVM)	Virtual SVM, deg-9 poly, 2-pixel jittered	None	Deskewing	0.56 <sup>[34]</sup>
Deep neural network (DNN)	2-layer 784-800-10	None	None	1.6 <sup>[35]</sup>
Deep neural network	2-layer 784-800-10	Elastic distortions	None	0.7 <sup>[35]</sup>
Deep neural network	6-layer 784-2500-2000-1500-1000-500-10	Elastic distortions	None	0.35 <sup>[36]</sup>
Convolutional neural network (CNN)	6-layer 784-40-80-500-1000-2000-10	None	Expansion of the training data	0.31 <sup>[37]</sup>
Convolutional neural network	6-layer 784-50-100-500-1000-10-10	None	Expansion of the training data	0.27 <sup>[38]</sup>
Convolutional neural network (CNN)	13-layer 64-128(5x)-256(3x)-512-2048-256-256-10	None	None	0.25 <sup>[22]</sup>
Convolutional neural network	Committee of 35 CNNs, 1-20-P-40-P-150-10	Elastic distortions	Width normalizations	0.23 <sup>[17]</sup>
Convolutional neural network	Committee of 5 CNNs, 6-layer 784-50-100-500-1000-10-10	None	Expansion of the training data	0.21 <sup>[24][25]</sup>
Random Multimodel Deep Learning (RMDL)	10 NN-10 RNN - 10 CNN	None	None	0.18 <sup>[27]</sup>
Convolutional neural network	Committee of 20 CNNs with Squeeze-and-Excitation Networks <sup>[39]</sup>	None	Data augmentation	0.17 <sup>[40]</sup>
Convolutional neural network	Ensemble of 3 CNNs with varying kernel sizes	None	Data augmentation consisting of rotation and translation	0.09 <sup>[41]</sup>

*What's the problem with this?*


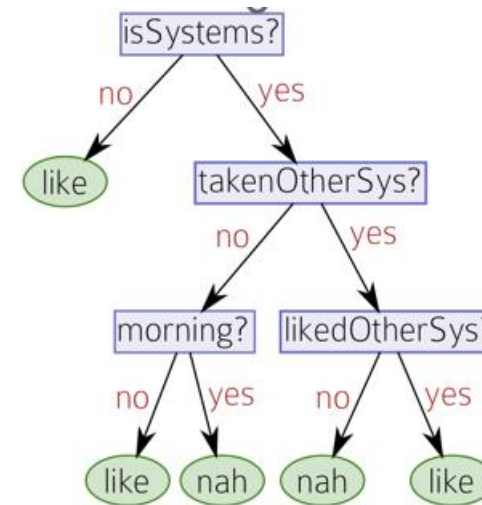
# Supervised learning setup



- Goal: design learning algorithm  $\mathcal{A}$  such that its output  $\hat{f}$  on iid training data  $S$  has low generalization error

# Terminologies

- Model: the predictor  $\hat{f}$ 
  - Often from a model class  $\mathcal{F}$ ,
  - e.g.  $\mathcal{F} = \{\text{decision trees}\}, \{\text{linear classifiers}\}$
- Parameter: specifics of  $\hat{f}$ 
  - E.g. for decision tree  $\hat{f}$ : tree structure, questions in nodes, labels in leaves
  - For linear classifier: linear coefficients
- Hyperparameter: specifics of learning algorithm  $\mathcal{A}$ 
  - E.g. in [DecisionTreeTrain](#), constrain to output tree of depth  $\leq h$
  - Tuning hyperparameters often results in {over, under}-fitting








# Hyperparameter tuning using validation set

- E.g. in decision tree training, how to choose tree depth  $h \in \{1, \dots, H\}$ ?

- For each hyperparameter  $h \in \{1, \dots, H\}$ :
  - Train  $\text{Tree}_h$  using [DecisionTreeTrain](#) by constraining the tree depth to be  $h$
- Choose one from  $\text{Tree}_1, \dots, \text{Tree}_H$



- Idea 1: choose  $\text{Tree}_h$  that minimizes training error 
- Idea 2: choose  $\text{Tree}_h$  that minimizes test error 
- Idea 3: further split training set to training set and validation set (development/hold-out set), (1) train  $\text{Tree}_h$ 's using the (new) training set; (2) choose  $\text{Tree}_h$  that minimizes validation error

Training: 700 examples

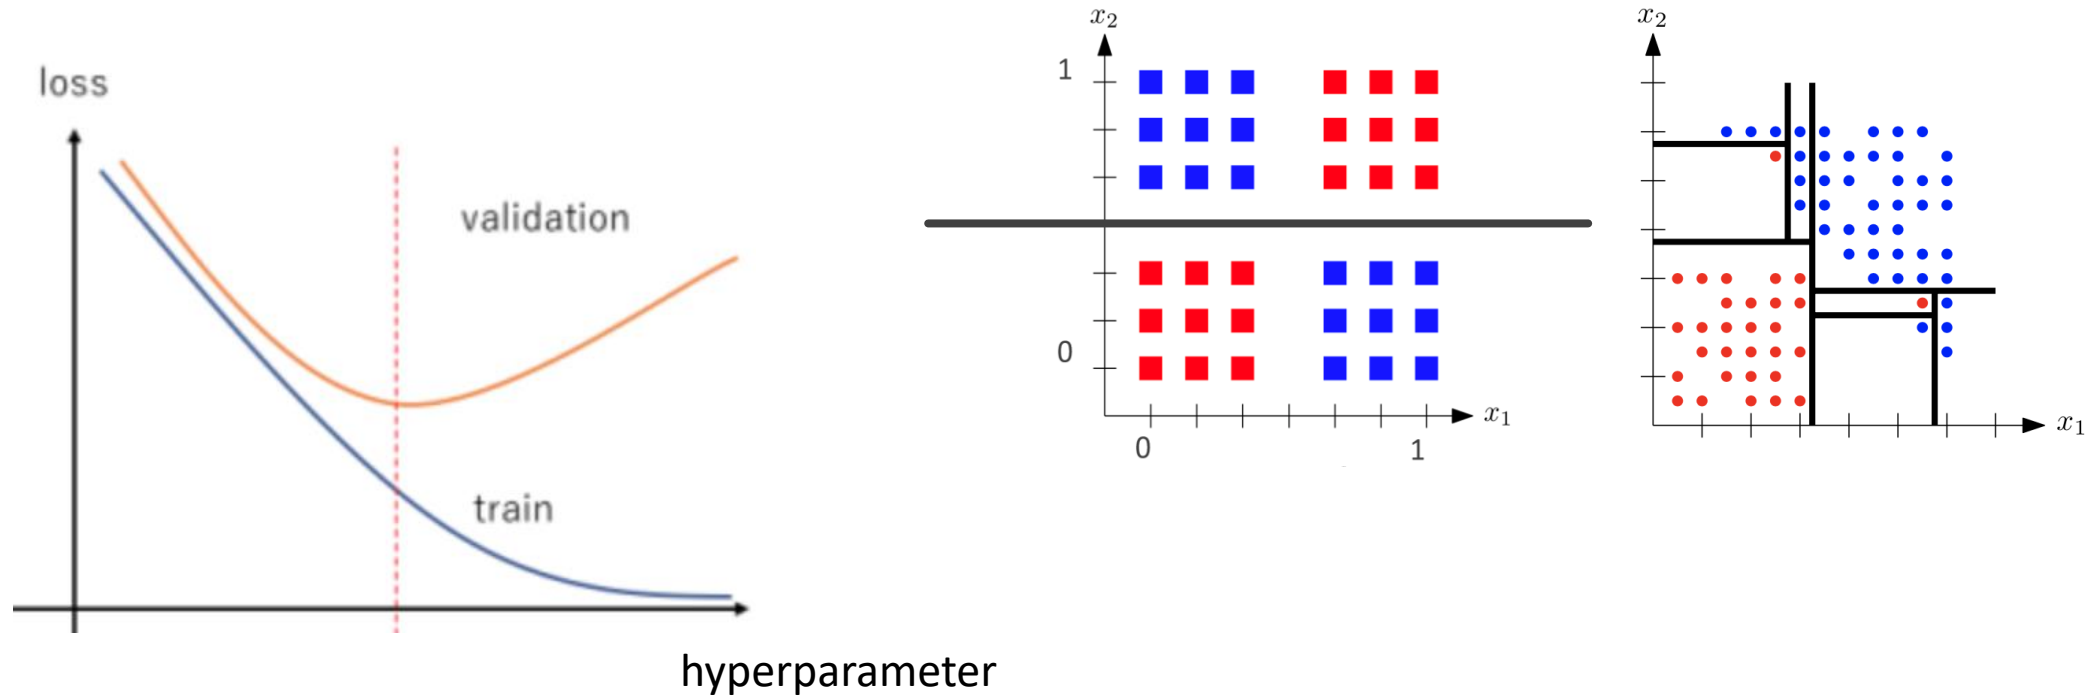
Val: 100  
examples

Test: 200  
examples



# Hyperparameter tuning using validation set


- E.g. in decision tree training, how to choose tree depth  $h \in \{1, \dots, H\}$ ?



- Law of large numbers  $\Rightarrow$  Validation error closely approximates generalization error (& test error)

# Overfitting vs Underfitting

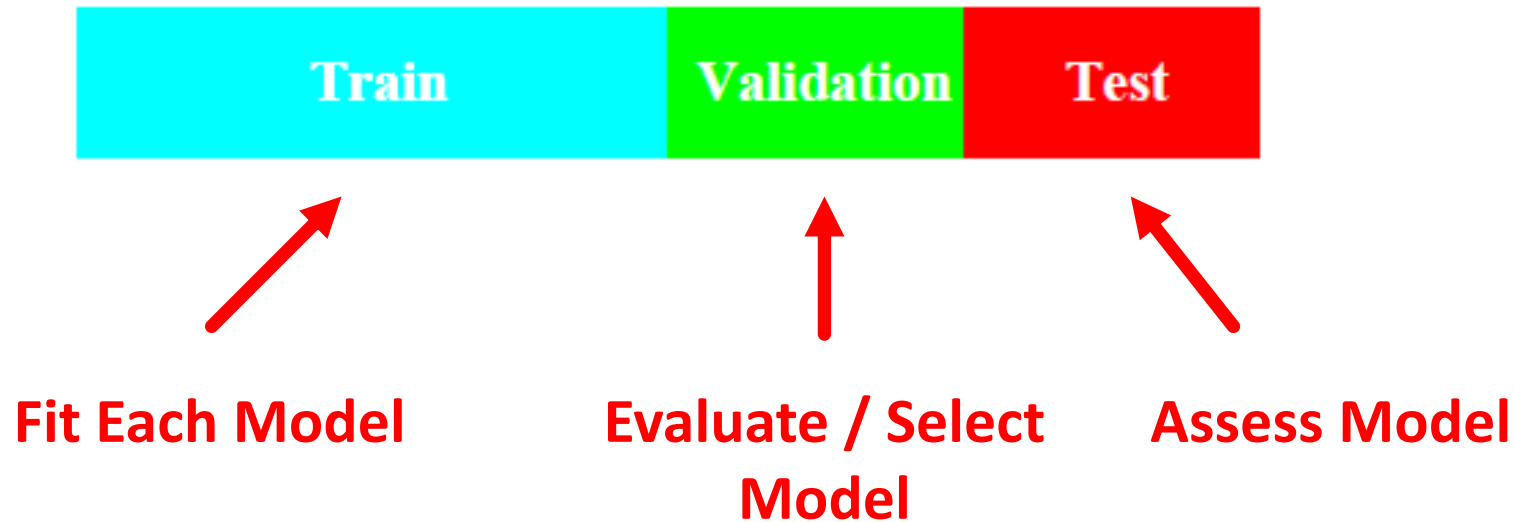
Underfitting: performs poorly on *both* training and validation...



...overfitting: performs well on training but not on validation

# Model Selection / Assessment


Partition your data into Train-Validation-Test sets



- Ideally, Test set is kept in a “vault” and only peek at it once model is selected
- Training-Validation-Test splits work if you have enough data (“data rich”)
- As a general rule 50% Training, 25% Validation, 25% Test (very loose rule)

# Hyperparameter tuning: cross-validation

**Main idea:** reuse data by splitting the training / validation data in multiple ways



**N-fold Cross Validation:** Partition training data into  $N$  “chunks” and for each run select one chunk to be validation data

For each run, fit to training data ( $N-1$  chunks) and measure accuracy on validation set. Average model error across all runs.

**Drawback** Need a lot of training data to partition.

# Cross-validation: formal description

- For hyperparameter  $h \in \{1, \dots, H\}$ 
  - For  $k \in \{1, \dots, K\}$ 
    - train  $\hat{f}_k^h$  with  $S \setminus \text{fold}_k$
    - measure error rate  $e_{h,k}$  of  $\hat{f}_k^h$  on  $\text{fold}_k$
  - Compute the average error of the above:  $\widehat{\text{err}}_h = \frac{1}{K} \sum_{k=1}^K e_{h,k}$
- Choose  $\hat{h} = \arg \min_h \widehat{\text{err}}_h$
- Train  $\hat{f}$  using  $S$  (all the training points) with hyperparameter  $\hat{h}$
- $k = |S|$ : leave one out cross validation (LOOCV)



# Inductive bias

- What classification problem is class A vs. class B?
  - Birds vs. Non-birds
  - Flying animals vs. non-flying animals



class A



class B

- **Inductive bias**: in the absence of data that narrow down the target concept, what type of solutions are we likely to prefer?
- What is the inductive bias of learning shallow decision trees?

# An example real-world machine learning pipeline

- Any step can go wrong
  - E.g. data collection, data representation
- Debugging pipeline: run *oracle experiments*
  - Assuming all lower-level tasks are perfectly done, is this step achieving what we want?
- General suggestions:
  - Build the stupidest thing that could possibly work
  - Decide whether / where to fix it

1	real world goal	increase revenue
2	real world mechanism	better ad display
3	learning problem	classify click-through
4	data collection	interaction w/ current system
5	collected data	query, ad, click
6	data representation	bow <sup>2</sup> , $\pm$ click
7	select model family	decision trees, depth 20
8	select training data	subset from april'16
9	train model & hyperparams	final decision tree
10	predict on test data	subset from may'16
11	evaluate error	zero/one loss for $\pm$ click
12	deploy!	(hope we achieve our goal)



# Next lecture (1/23)

- Geometric view of supervised learning; nearest neighbor methods
- Assigned reading: CIML Chap. 3 (Geometry and Nearest Neighbors)
- HW1

# Simple case: discrete domain $\mathcal{X}$


$P_D(x, y)$	$x = 1$	$x = 2$	$x = 3$
$y = -1$	0.2	0.2	0.15
$y = +1$	0.1	0.3	0.05

Which classifier is better?

- $f_1(1) = -1, f_1(2) = -1, f_1(3) = -1 \Rightarrow L_D(f_1) = 0.1 + 0.3 + 0.05$
- $f_2(1) = -1, f_2(2) = +1, f_2(3) = -1 \Rightarrow L_D(f_2) = 0.1 + 0.2 + 0.05$
- What is the best classifier?
- For any  $x$ , should choose  $y$  that has higher value of  $P_D(x, y)$
- $f^*(1) = -1, f^*(2) = +1, f^*(3) = -1$

# Bayes optimal classifier

- $f_{BO}(x) = \arg \max_{y \in \mathcal{Y}} P_D(X = x, Y = y) = \arg \max_{y \in \mathcal{Y}} P_D(Y = y | X = x), \forall x \in \mathcal{X}$
- Theorem:  $f_{BO}$  achieves the smallest error rate among all functions.
- Bayes error rate:  $L_D(f_{BO})$



# Proof of theorem

- Step 1: consider accuracy:

- $A_D(f) = 1 - L_D(f) = P_D(Y = f(X)) = \sum_x P_D(X = x, Y = f(x))$
- Suffices to show  $f_{BO}$  has the highest accuracy

- Step 2: comparison:

$$A_D(f_{BO}) - A_D(f) = \sum_x P_D(X = x, Y = f_{BO}(x)) - P_D(X = x, Y = f(x)) \geq 0$$


$$f_{BO}(x) = \arg \max_{y \in \mathcal{Y}} P_D(X = x, Y = y)$$

- Remark: similar reasoning can be used to prove the theorem with continuous domain  $\mathcal{X}$  (sum  $\rightarrow$  integral)

# When is the Bayes error rate nonzero?



- $L_D(f_{BO}) = \sum_x \min( P_D(Y = +1, X = x), P_D(Y = -1, X = x))$

- Limited feature representation



- feature noise
  - Sensor failure
  - Typo in reviews for sentiment classification

- label noise
  - Crowdsourcing settings



# Class Participation

- Asking review questions on Piazza (3pts)
  - Every week, I will ask two of you to post questions (related to the past week's material) on Piazza
  - 3 questions per student
- Other in-class / Piazza discussions (e.g. asking/answering in-class questions; Piazza Q&As)
- Extra credit: Catching errors in the CIML book
  - Post on Piazza; we'll discuss and confirm together, and hopefully send these back to the author
  - 1pt for every error found