

CSC 665: Homework 3

Chicheng Zhang

October 31, 2019

Please complete the following set of problems. You must do the exercises completely on your own (no collaboration allowed). The exam is due **on Nov 14, 12:30pm, on Gradescope**. You are free to cite existing theorems from the textbooks and course notes.

Problem 1

Consider the homogeneous, soft-margin SVM optimization problem:

$$\underset{w, \xi}{\text{minimize}} \quad \frac{\lambda}{2} \|w\|^2 + \sum_{i=1}^n \xi_i \tag{1}$$

$$\begin{aligned} \text{s. t.} \quad & y_i (\langle w, x_i \rangle) \geq 1, & \forall i \in \{1, \dots, n\}, \\ & \xi_i \geq 0, & \forall i \in \{1, \dots, n\}. \end{aligned} \tag{2}$$

1. Introducing dual variables $\alpha_i \geq 0$ for each constraint i , $i \in \{1, \dots, n\}$ and $\beta_i \geq 0$ for each constraint i , $i \in \{1, \dots, n\}$, compute the Lagrangian function $L(w, \xi, \alpha, \beta)$.
2. Derive the dual optimization problem.
3. Use the KKT condition to interpret: which of the training examples are “support vectors” that contribute to the SVM solution?

Problem 2

Suppose we have k finite hypothesis classes $\mathcal{H}_1, \dots, \mathcal{H}_k$, and m training examples drawn iid from D . In addition we are given the promise that there exists $i_0 \in \{1, \dots, k\}$ such that $\min_{h \in \mathcal{H}_{i_0}} \text{err}(h, D) = 0$ (but we don't know the identity of i_0); Can we design an algorithm that has generalization error $O(\frac{\ln |\mathcal{H}_{i_0}|}{m})$ with high probability? Why or why not?

Problem 3

Show that for AdaBoost, at iteration t , the updated distribution D_{t+1} satisfies that

$$\sum_{i=1}^m D_{t+1}(i) \mathbf{1}(h_t(x_i) \neq y_i) = \frac{1}{2}.$$

Why is this a reasonable update?

Problem 4

In this exercise, we conduct experiment with AdaBoost with a simple benchmark dataset named *ringnorm*.

1. Generate 100 training and 100 test examples from the following distribution D supported on $\mathbb{R}^{10} \times \{\pm 1\}$: $\mathbb{P}_D(Y = -1) = \mathbb{P}_D(Y = +1) = \frac{1}{2}$, $X|_{Y=+1} \sim N((0, \dots, 0), 4I)$; $X|_{Y=-1} \sim N((\frac{2}{\sqrt{20}}, \dots, \frac{2}{\sqrt{20}}), I)$.
2. Define base hypothesis class $\mathcal{H} = \{\sigma \cdot (2I(x_i \leq t) - 1), \sigma \in \{\pm 1\}, i \in \{1, \dots, d\}, t \in \mathbb{R}\}$ as the set of bi-directional decision stumps. Let the weak learner \mathcal{B} be: given a weighted dataset, return the classifier $h \in \mathcal{H}$ that has the smallest weighted error. Implement AdaBoost with \mathcal{B} , and run it for 300 iterations. At time t , suppose the following cumulative voting classifier

$$H_t(x) = \text{sign}(f_t(x)), \quad f_t(x) = \sum_{s=1}^t \alpha_s h_s(x)$$

is produced.

Plot AdaBoost's learning curves: the training error of H_t , the test error of H_t and the training exponential loss of f_t as a function of iteration t . What do you see?

3. Given a voting classifier f_t , define its normalization as

$$\bar{f}_t(x) = \frac{f_t(x)}{\sum_{s=1}^t \alpha_s} = \frac{\sum_{s=1}^t \alpha_s h_s(x)}{\sum_{s=1}^t \alpha_s}. \quad (3)$$

Now, given an example (x, y) , define its normalized margin at timestep t as $y\bar{f}_t(x)$. At iterations $t = 10, 30, 50, 100, 300$, show histograms of normalized margins of training examples. Do you see any tendency at t increases?

Problem 5 (No need to submit)

Show that AdaBoost produces large-margin voting classifiers under the γ -weak learning assumption. If at every iteration t , $\epsilon_t \leq \frac{1}{2} - \gamma$, then after T rounds, the *margin error* of the output classifier will also decrease exponentially in T . Specifically, show:

$$\frac{1}{m} \sum_{i=1}^m \mathbf{1}(y_i \bar{f}_T(x_i) \leq \gamma) \leq \exp\{-\Omega(T\gamma^2)\}.$$

where \bar{f}_T is the normalized voting classifier defined as per Equation (3).